

# RentalListingVisualizations

April 4, 2017

## 1 Initializing Dataset

We must initialize the dataset first, loading it in as a pandas object from a JSON object.

```
In [43]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
colors = sns.color_palette()
```

```
trainData = pd.read_json("train.json")
testData = pd.read_json("test.json")
trainData.head()
```

```
Out[43]:
```

	bathrooms	bedrooms	building_id \
10	1.5	3	53a5b119ba8f7b61d4e010512e0dfc85
10000	1.0	2	c5c8a357cba207596b04d1afd1e4f130
100004	1.0	1	c3ba40552e2120b0acfc3cb5730bb2aa
100007	1.0	1	28d9ad350afeaab8027513a3e52ac8d5
100013	1.0	4	0

	created \
10	2016-06-24 07:54:24
10000	2016-06-12 12:19:27
100004	2016-04-17 03:26:41
100007	2016-04-18 02:22:02
100013	2016-04-28 01:32:41

	description \
10	A Brand New 3 Bedroom 1.5 bath ApartmentEnjoy ...
10000	
100004	Top Top West Village location, beautiful Pre-w...
100007	Building Amenities - Garage - Garden - fitness...
100013	Beautifully renovated 3 bedroom flex 4 bedroom...

	display_address \
--	-------------------

10	Metropolitan Avenue				
10000	Columbus Avenue				
100004	W 13 Street				
100007	East 49th Street				
100013	West 143rd Street				

		features	interest_level
10		[]	medium
10000	[Doorman, Elevator, Fitness Center, Cats Allow...		low
100004	[Laundry In Building, Dishwasher, Hardwood Flo...		high
100007	[Hardwood Floors, No Fee]		low
100013	[Pre-War]		low

	latitude	listing_id	longitude	manager_id
10	40.7145	7211212	-73.9425	5ba989232d0489da1b5f2c45f6688adc
10000	40.7947	7150865	-73.9667	7533621a882f71e25173b27e3139d83d
100004	40.7388	6887163	-74.0018	d9039c43983f6e564b1482b273bd7b01
100007	40.7539	6888711	-73.9677	1067e078446a7897d2da493d2f741316
100013	40.8241	6934781	-73.9493	98e13ad4b495b9613cef886d79a6291f

		photos	price	\
10	[https://photos.renthop.com/2/7211212_1ed4542e...		3000	
10000	[https://photos.renthop.com/2/7150865_be3306c5...		5465	
100004	[https://photos.renthop.com/2/6887163_de85c427...		2850	
100007	[https://photos.renthop.com/2/6888711_6e660cee...		3275	
100013	[https://photos.renthop.com/2/6934781_1fa4b41a...		3350	

	street_address
10	792 Metropolitan Avenue
10000	808 Columbus Avenue
100004	241 W 13 Street
100007	333 East 49th Street
100013	500 West 143rd Street

In [44]: testData.head()

Out [44]:

	bathrooms	bedrooms	building_id	\
0	1.0	1	79780be1514f645d7e6be99a3de696c5	
1	1.0	2		0
100	1.0	1	3dbbb69fd52e0d25131aa1cd459c87eb	
1000	1.0	2	783d21d013a7e655bddc4ed0d461cc5e	
100000	2.0	2	6134e7c4dd1a98d9aee36623c9872b49	

	created	\
0	2016-06-11 05:29:41	
1	2016-06-24 06:36:34	
100	2016-06-03 04:29:40	
1000	2016-06-11 06:17:35	

100000 2016-04-12 05:24:17

	description \
0	Large with awesome terrace--accessible via bed...
1	Prime Soho - between Bleecker and Houston - Ne...
100	New York chic has reached a new level ...
1000	Step into this fantastic new Construction in t...
100000	~Take a stroll in Central Park, enjoy the ente...

	display_address \
0	Suffolk Street
1	Thompson Street
100	101 East 10th Street
1000	South Third Street\r
100000	Midtown West, 8th Ave

	features	latitude \
0	[Elevator, Laundry in Building, Laundry in Uni...	40.7185
1	[Pre-War, Dogs Allowed, Cats Allowed]	40.7278
100	[Doorman, Elevator, No Fee]	40.7306
1000	[Roof Deck, Balcony, Elevator, Laundry in Buil...	40.7109
100000	[Common Outdoor Space, Cats Allowed, Dogs Allo...	40.7650

	listing_id	longitude	manager_id \
0	7142618	-73.9865	b1b1852c416d78d7765d746cb1b8921f
1	7210040	-74.0000	d0b5648017832b2427eeb9956d966a14
100	7103890	-73.9890	9ca6f3baa475c37a3b3521a394d65467
1000	7143442	-73.9571	0b9d5db96db8472d7aeb67c67338c4d2
100000	6860601	-73.9845	b5eda0eb31b042ce2124fd9e9fcfce2f

	photos	price \
0	[https://photos.renthop.com/2/7142618_1c45a2c8...	2950
1	[https://photos.renthop.com/2/7210040_d824cc71...	2850
100	[https://photos.renthop.com/2/7103890_85b33077...	3758
1000	[https://photos.renthop.com/2/7143442_0879e9e0...	3300
100000	[https://photos.renthop.com/2/6860601_c96164d8...	4900

	street_address
0	99 Suffolk Street
1	176 Thompson Street
100	101 East 10th Street
1000	251 South Third Street\r
100000	260 West 54th Street

Notice that the training set and the test set have the same factors besides interest\_level. That's because interest\_level is our target variable, and we'll use it for comparison and analysis in the following below.


## 2 Visualizing Data

We want to start taking a look at what variables have most influence over the target variable. So, we want to compare the distributions of the varying factors.

### 2.0.1 Interest Levels

```
In [45]: interestLevelCounts = trainData['interest_level'].value_counts()

plt.figure(figsize=(16,8))
sns.barplot(interestLevelCounts.index, interestLevelCounts.values, alpha=0.5)
plt.ylabel('Number of Occurrences', fontsize=12)
plt.xlabel('Interest level', fontsize=12)
plt.show()
```



The amount of 'low' interest level houses outnumbers the amount of 'medium' and 'high' occurrences. This means the distribution is a bit unbalanced, and we need to be careful when analyzing direct numbers of the factors. We may want to consider using percentages when comparing variables.


### 2.0.2 Bathrooms

```
In [46]: # Normalizing the outliers
outlierBathroomRows = trainData['bathrooms'] > 4
trainData.loc[outlierBathroomRows, 'bathrooms'] = 4

bathroomCount = trainData['bathrooms'].value_counts()


plt.figure(figsize=(16,8))
sns.barplot(bathroomCount.index, bathroomCount.values, color=colors[3])
```

```
plt.ylabel('Number of Occurances', fontsize=12)
plt.xlabel('Number of Bathrooms', fontsize=12)
plt.show()
```



```
In [47]: #trainData['bathrooms'].ix[trainData['bathrooms'] > 4] = 4
```

```
plt.figure(figsize=(16,8))
sns.violinplot(trainData['bathrooms'], trainData['interest_level'], order=
plt.xlabel('Number of Bathrooms', fontsize=12)
plt.ylabel('Interest Levels', fontsize=12)
plt.show()
```



The number of bathrooms appears to be pretty consistent across interest levels. So bathrooms will not be a very important factor in analysis for the target variable, at least at first glance.


### 2.0.3 Prices

Initially after graphing the prices there were a couple of outliers throwing off the visualization, so we left that out in favor of a more useful approach.

```
In [48]: # Normalizing the outliers
percentile_99_Price = np.percentile(trainData['price'], 99)
outlierPriceRows = trainData['price'] > percentile_99_Price
trainData.loc[outlierPriceRows, 'price'] = percentile_99_Price

plt.figure(figsize=(16,8))
sns.distplot(trainData['price'])
plt.xlabel('Price values', fontsize=12)
plt.ylabel('Density', fontsize=12)
plt.show()
```


```
/Users/JustinonTG/anaconda/lib/python3.6/site-packages/statsmodels/nonparametric/kde
y = X[:m/2+1] + np.r_[0,X[m/2+1:],0]*1j
```




That overall distribution looks quite nice, with a skew right. It should be interesting to visualize the prices against the interest levels.

```
In [49]: plt.figure(figsize=(16,8))
sns.violinplot(x='interest_level', y='price', data=trainData, order=['low
```

```
plt.xlabel('Interest Level', fontsize=12)
plt.ylabel('Price values', fontsize=12)
plt.show()
```




```
In [50]: plt.figure(figsize=(16,8))
sns.boxplot(x='price', y='interest_level', data=trainData, order=['low', 'medium', 'high'])
plt.xlabel('Prices', fontsize=12)
plt.ylabel('Interest Levels', fontsize=12)
plt.show()
```



It's easy to notice here that the 'low' interest level flats also are the ones that have the most density in the higher price values. Also, the 'high' interest level flats have a lower average price. This is logical, and helps in our prediction. The price variable therefore should have some more weight than the other factors in predicting the influence level.

## 2.0.4 Bedrooms


```
In [51]: plt.figure(figsize=(16,8))
sns.violinplot(x='interest_level', y='bedrooms', data=trainData, order=["low", "medium", "high"])
plt.xlabel('Interest Level')
plt.ylabel('Number of Bedrooms')
plt.show()
```



```
In [62]: # Normalizing the outliers
percentile_995_Bedrooms = np.percentile(trainData['bedrooms'], 99.5)
percentile_005_Bedrooms = np.percentile(trainData['bedrooms'], .5)
outlierHighBedRows = trainData['bedrooms'] > percentile_995_Bedrooms
outlierLowBedRows = trainData['bedrooms'] < percentile_005_Bedrooms
trainData.loc[outlierHighBedRows, 'bedrooms'] = percentile_995_Bedrooms
trainData.loc[outlierLowBedRows, 'bedrooms'] = percentile_005_Bedrooms

plt.figure(figsize=(16,8))
sns.violinplot(x='bedrooms', y='interest_level', data=trainData, order=["0", "1", "2", "3", "4", "5", "6", "7", "8"])
plt.xlabel('Number of Bedrooms')
plt.ylabel('Interest Levels')
plt.show()
```






### 2.0.5 Features

Here we are going to display the various numbers of features and how they're distributed according to the interest level. This will be achieved by counting the number of features provided for each listing and displaying the counts.

```
In [53]: trainData['num_features'] = trainData['features'].apply(len)
```


```
plt.figure(figsize=(16,8))
listOrder = ['low', 'medium', 'high']
sns.boxplot(x='num_features', y='interest_level', data=trainData, order=listOrder)
plt.xlabel('Number of Features', fontsize=12)
plt.ylabel('Interest Level')
plt.show()
```



Normalizing the outliers yields us the following:

```
In [54]: percentile_995_Features = np.percentile(trainData['num_features'], 99.5)
percentile_005_Features = np.percentile(trainData['num_features'], .5)
outlierHighFeatRows = trainData['num_features'] > percentile_995_Features
outlierLowFeatRows = trainData['num_features'] < percentile_005_Features
trainData.loc[outlierHighFeatRows, 'num_features'] = percentile_995_Features
trainData.loc[outlierLowFeatRows, 'num_features'] = percentile_005_Features

plt.figure(figsize=(16,8))
listOrder = ['low', 'medium', 'high']
sns.violinplot(x='num_features', y='interest_level', data=trainData, order=listOrder)
plt.xlabel('Number of Features', fontsize=12)
plt.ylabel('Interest Level')
plt.show()
```




## 2.1 Geolocation

### 2.1.1 Latitude

```
In [55]: # Normalizing the outliers
percentile_995_Latitude = np.percentile(trainData['latitude'], 99.5)
percentile_005_Latitude = np.percentile(trainData['latitude'], .5)
outlierUpperLatRows = trainData['latitude'] > percentile_995_Latitude
outlierLowerLatRows = trainData['latitude'] < percentile_005_Latitude
trainData.loc[outlierUpperLatRows, 'latitude'] = percentile_995_Latitude
trainData.loc[outlierLowerLatRows, 'latitude'] = percentile_005_Latitude


plt.figure(figsize=(16,8))
sns.distplot(trainData['latitude'], color=colors[5])
plt.xlabel("Latitude", fontsize=12)
plt.ylabel("Density", fontsize=12)
plt.show()

/Users/JustinonTG/anaconda/lib/python3.6/site-packages/statsmodels/nonparametric/kde
y = X[:m/2+1] + np.r_[0,X[m/2+1:],0]*1j
```



Interestingly, the latitude appears to follow a Gaussian distribution. Further bootstrap testing and Q-Q plotting should be conducted here to compare to the normal.

```
In [56]: plt.figure(figsize=(16,8))
sns.boxplot(x='latitude', y='interest_level', data=trainData, order=["low", "medium", "high"])
plt.xlabel("Latitude", fontsize=12)
plt.ylabel("Interest Levels", fontsize=12)
plt.show()
```




## 2.1.2 Longitude


```
In [57]: # Normalizing the outliers
percentile_9975_Longitude = np.percentile(trainData['longitude'], 99.75)
percentile_0025_Longitude = np.percentile(trainData['longitude'], .25)
outlierUpperLongRows = trainData['longitude'] > percentile_9975_Longitude
outlierLowerLongRows = trainData['longitude'] < percentile_0025_Longitude
trainData.loc[outlierUpperLongRows, 'longitude'] = percentile_9975_Longitude
trainData.loc[outlierLowerLongRows, 'longitude'] = percentile_0025_Longitude

plt.figure(figsize=(16,8))
sns.distplot(trainData['longitude'], color=colors[4])
plt.xlabel("Longitude", fontsize=12)
plt.ylabel("Density", fontsize=12)
plt.show()
```

```
/Users/JustinonTG/anaconda/lib/python3.6/site-packages/statsmodels/nonparametric/kde
y = X[:m/2+1] + np.r_[0,X[m/2+1:],0]*1j
```



```
In [58]: plt.figure(figsize=(16,8))
sns.boxplot(x='longitude', y='interest_level', data=trainData, order=["low
plt.xlabel("Longitude", fontsize=12)
plt.ylabel("Interest Levels", fontsize=12)
plt.show()
```




### 2.1.3 Mapping

First we'll plot the latitude and longitude together to graph the map of the city with which we're dealing.

```
In [59]: plt.figure(figsize=(16,16))
          sns.jointplot(x='longitude',y='latitude', data=trainData)
          plt.show()
```

<matplotlib.figure.Figure at 0x13d1772e8>

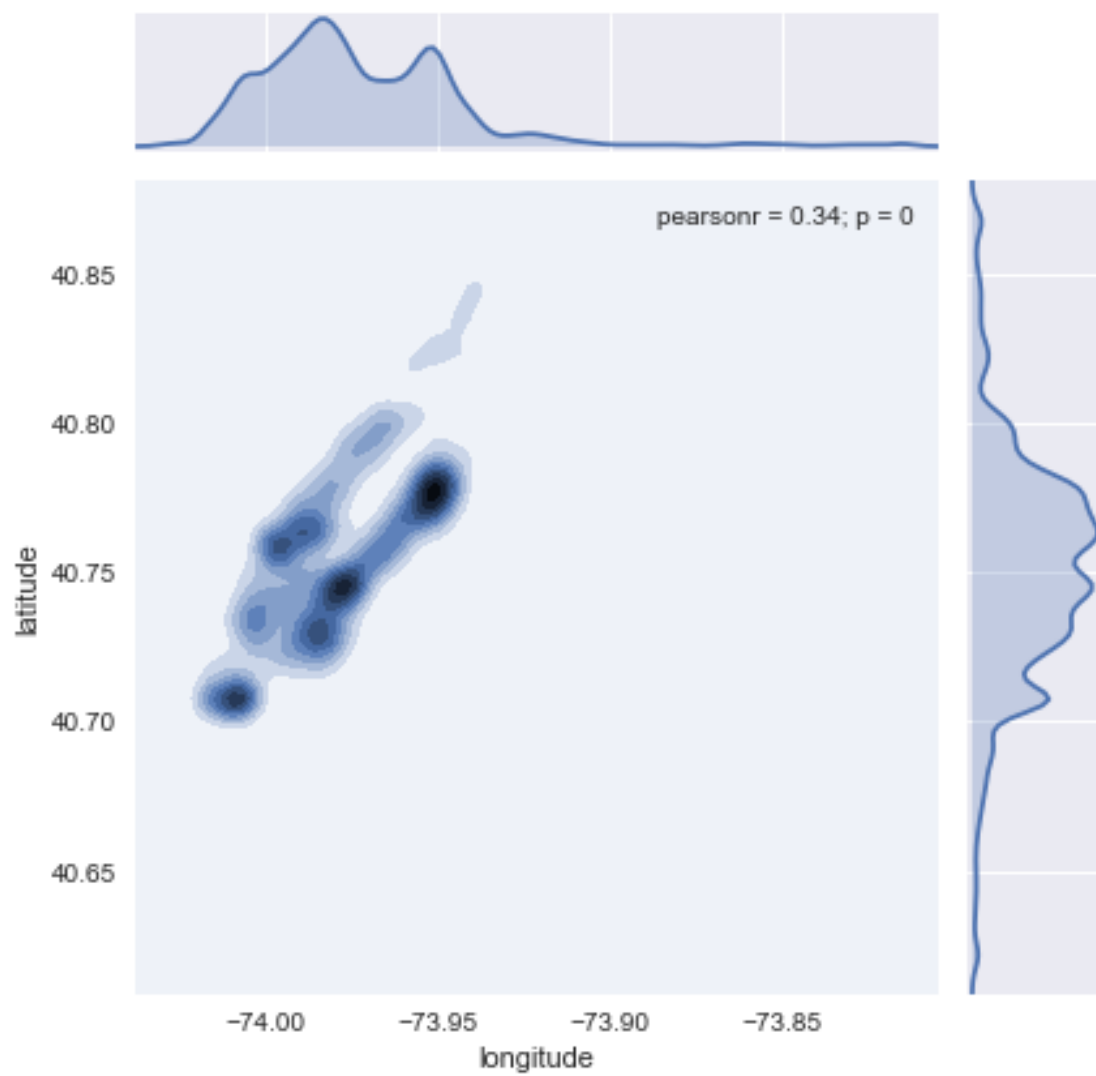


Now to check for the actual densities within the region, we'll use the kernel density estimation joint plot.

```
In [60]: plt.figure(figsize=(16,16))
          sns.jointplot(x='longitude',y='latitude', data=trainData, kind='kde')
          plt.show()
```

```
/Users/JustinonTG/anaconda/lib/python3.6/site-packages/statsmodels/nonparametric/kde
y = X[:m/2+1] + np.r_[0,X[m/2+1:],0]*1j
```

```
<matplotlib.figure.Figure at 0x131521630>
```



In [ ]: